

canjob

Intelligent Candidate Discovery & Ranking

A JD-faithful, CPU-only, fully offline ranking engine that scores 100,000 candidates in ~30 seconds and explains every pick with profile evidence.

Problem

Data & AI Challenge:
Intelligent Candidate Discovery

Team

canjob

Team Leader

Navpreet Singh

What we built and why it is different

- ▶ **Solution:** an ensemble ranker that turns a JD into config, then scores every candidate with three independent lenses and fuses them.
- ▶ **Reads profiles, not just keywords:** a deterministic JD-faithful rules model encodes must-haves and the disqualifiers the JD spells out.
- ▶ **Behavior-aware:** the 23 Redrob signals modulate score so unreachable perfect-on-paper profiles sink.
- ▶ **Honest by design:** logically-impossible honeypots are detected and removed; every row carries evidence-based reasoning.

3 lenses -> 1 rank

Rules + lexical TF-IDF + MiniLM semantic, fused with Reciprocal Rank Fusion.

No LLM at rank time

Differentiator vs GPT-per-candidate systems that cannot scale to a 200k pool in production.

Config-driven

One codebase serves many jobs; each JD is its own versioned config folder.

Key requirements and the signals that matter

Extracted from the JD

- ▶ Embeddings & dense/hybrid retrieval; vector DBs (Pinecone, FAISS, Qdrant)
- ▶ Ranking / recsys + evaluation literacy (NDCG, MRR, MAP, A/B)
- ▶ LLM fine-tuning / RAG, production Python at scale
- ▶ Product-company builder profile, 5-9 yrs, founding-team mindset

Beyond keyword matching

- ▶ **Semantic fit:** MiniLM matches meaning so a recommender-systems builder ranks even without the buzzwords.
- ▶ **Availability:** response rate, recency, open-to-work, interview/offer reliability, recruiter demand.
- ▶ **Trust:** endorsements + skill duration discount lazy keyword stuffing.
- ▶ **Disqualifiers:** off-domain title with stuffed AI skills, framework-only, title-chaser tenures.

Retrieve, score, rank, fuse

Rules (deterministic)

JD-faithful score from config: skill facets, experience band, product vs services, location. Minus explicit penalties. x behavioral modifier.

Lexical (TF-IDF)

Several focused queries (retrieval, ranking/eval, full JD) for keyword recall instead of one blurry mega-query.

Semantic (MiniLM)

Offline sentence embeddings, precomputed once. Cosine-to-JD captures meaning the other lenses miss.

Fusion: Reciprocal Rank Fusion (RRF)

Each lens produces a ranking; fused score = $\text{sum of } \text{weight} / (k + \text{rank})$. Combining on ranks, not raw scores, means no single lens scale can dominate. Weights live in jd_facets.json. Honeypots are excluded; ties break by candidate_id; scores are non-increasing.

Grounded reasoning, no hallucination, suspicious profiles

- ▶ **Evidence-based:** reasoning cites the candidate's own named skills (proficiency + months), employers, and evaluation terms found in their text.
- ▶ **Mapped to the JD:** each clause connects a concrete fact to a specific JD requirement, not generic praise.
- ▶ **Honest concerns:** gaps (services-only career, framework-only, long notice) are stated, and tone matches the rank.
- ▶ **No hallucination:** every claim is pulled from the profile; nothing is invented, so Stage-4 checks pass.


0 in top 100

Honeypots (impossible profiles: expert-but-0-months, tenure > experience, role longer than career) are detected and removed. Threshold for disqualification is >10%.

High precision

We only flag logically-impossible profiles, not noisy heuristics, so legitimate candidates are never wrongly dropped.

From JD to ranked CSV



Offline boundary: only Precompute may use the network (one-time model download) and may exceed 5 min.

The Rank step that produces the CSV is CPU-only, offline, and finishes in ~30s for the full 100k pool.

Components and data flow

INPUTS (offline)

candidates.jsonl

100k profiles, each with 23 behavioral signals; managed via the web UI.

per-job config

JD facets, disqualifier filters and fusion weights; one folder per JD.

semantic_scores.npz

0.6 MB MiniLM cosine-to-JD scores, precomputed once and committed.

RANKING ENGINE (CPU, ~30s)

featurize

one vectorized pandas/numpy pass: profile features + 23-signal availability modifier; flags impossible honeypots.

3 scorers

rules (JD-faithful) | TF-IDF (lexical) | MiniLM (semantic).

RRF fusion

rank-based merge of the three scorers; honeypots excluded.

OUTPUT + ACCESS

top-100 CSV

candidate_id, rank, 0-1 fit score, and evidence-based reasoning. Validator-clean, UTF-8.

run it anywhere

rank.py one command | Docker (--network none) | Colab sandbox | FastAPI web UI to upload JDs, manage candidates, and run matches.

Data flow: profiles + JD config -> one vectorized featurization (+ honeypot flags) -> three scorers -> RRF -> honeypot-filtered top 100 with evidence.

Quality and constraint compliance

~30s

full 100k ranked on CPU
(Docker ~50s)

0

honeypots in the submitted top
100

0.6 MB

committed semantic artifact
(git-friendly)

offline

verified with docker --network
none

Constraint compliance

- ▶ Runtime ≤ 5 min: rank step ~30s for 100k
- ▶ Memory ≤ 16 GB: single vectorized pass
- ▶ CPU only / network off: no GPU, no API calls
- ▶ Disk ≤ 5 GB: artifact 0.6 MB, optional cache ~150 MB

Confidence proxy (no ground truth)

- ▶ Scored on NDCG@10 (.50), NDCG@50 (.30), MAP (.15), P@10 (.05)
- ▶ Independent lenses cross-validate: a strong consensus core ranks in every strategy
- ▶ Agreement is reported by the ranker as a defensibility signal

Drive the same engine from a browser

candidate → job · intelligent discovery & ranking

[Jobs](#)
[Candidates](#)
[Matching](#)

Pool

100,000

Run match...

Filter by job

All jobs

Filter by candidate set

All candidate sets

Runs (3)

Job	Candidate set	Status	Pool	Top K	Took	Actions
Senior AI Engineer (Founding Team) - Redrob AI	Sample (first 50)	done	50	100	10s	View CSV PDF Re-run
Senior AI Engineer (Founding Team) - Redrob AI	Senior (8y+ experience)	done	400	100	7s	View CSV PDF Re-run
Senior AI Engineer (Founding Team) - Redrob AI	AI / ML engineers	done	400	100	4s	View CSV PDF Re-run

Results

Senior AI Engineer (Founding Team) - Redrob AI × AI / ML engineers · 100 ranked · 400 in pool · took 4s

[Export CSV](#)
[Export PDF](#)
[Close](#)

#	Candidate	Score	Reasoning
1	CAND_0018499	1.0000	Senior Machine Learning Engineer, 7.2y @ Zomato, Google. JD match -> vector DBs/hybrid search: Weaviate(exp,88mo), Pinecone(adv,54mo), Milvus(exp,40mo); embeddings retrieval: Embeddings(exp,65mo); ranking/recsys: Recommendation Systems(exp,46mo), Information Retrieval(adv,27mo), Learning to Rank(exp,70mo); eval frameworks: NDCG, MRR, A/B tests. Signals: 61% recruiter response, active 2026-05, 15d notice.
2	CAND_0046525	0.9634	Senior Machine Learning Engineer, 6.1y @ Genpact AI, LinkedIn. JD match -> vector DBs/hybrid search: Elasticsearch(adv,31mo), pgvector(adv,20mo), Qdrant(exp,44mo); embeddings retrieval: Sentence Transformers(exp,74mo); ranking/recsys: Information Retrieval(adv,59mo); eval frameworks: NDCG, MRR, A/B tests. Signals: 88% recruiter response, active 2026-05, 60d notice.
3	CAND_0011687	0.9386	Senior Nlp Engineer, 7.8y @ Niramai, Krutrim. JD match -> vector DBs/hybrid search: OpenSearch(exp,92mo), FAISS(adv,42mo), Weaviate(exp,60mo); embeddings retrieval: Embeddings(exp,56mo), Semantic Search(adv,45mo); eval frameworks: NDCG, MRR, A/B tests; LLM fine-tuning: PEFT(exp,75mo). Signals: 89% recruiter response, active 2026-05, 15d notice.
4	CAND_0002025	0.9321	Senior AI Engineer, 5.9y @ Apple, Aganitha. JD match -> vector DBs/hybrid search: FAISS(exp,86mo), OpenSearch(adv,50mo), Haystack(exp,64mo); embeddings retrieval: Sentence Transformers(exp,81mo); ranking/recsys: Recommendation Systems(exp,56mo); eval frameworks: A/B tests, offline eval. Signals: 80% recruiter response, active 2026-05, 30d notice.
5	CAND_0046064	0.8917	Senior Nlp Engineer, 8.9y @ Salesforce, Verloop.io. JD match -> vector DBs/hybrid search: Pinecone(exp,85mo), Haystack(exp,75mo), OpenSearch(adv,21mo); embeddings retrieval: Semantic Search(adv,43mo); ranking/recsys: BM25(exp,64mo); eval frameworks: NDCG, MRR, A/B tests. Signals: 78% recruiter response, active 2026-04, 30d notice.
6	CAND_0017960	0.8913	Recommendation Systems Engineer, 7.7y @ Nykaa, Wysa. JD match -> vector DBs/hybrid search: Haystack(exp,75mo), Qdrant(exp,45mo); ranking/recsys: Information Retrieval(exp,67mo); eval frameworks: A/B tests, offline eval; LLM fine-tuning: QLoRA(exp,63mo), PEFT(exp,71mo). Signals: 72% recruiter response, active 2026-04, 60d notice.

► **Jobs & candidate sets:** default JD pre-loaded (markdown); curate candidate sets from the 100k pool.

► **Run-match popup:** tick job(s) x candidate set(s); queues every combination with live progress.

► **Filterable runs + paging:** filter runs by job or set; results paginate 100 per page.

► **Import / export:** CSV in (validated) and CSV or print-ready PDF out, in one click.

► **Same engine:** identical featurize -> 3 lenses -> RRF -> honeypot filter as rank.py.

100k pool ranked end-to-end, 0 honeypots in the top 100.

Independent proof the ranking works

HuggingFace resume-job-description-fit: 1,759 human-labelled resume x JD pairs (Good / Potential / No-Fit), 71 jobs, 477 resumes. For each job we rank the WHOLE pool of 477 and see how high the recruiter's Good-Fit picks land.

On data we never tuned on, CanJob surfaces good-fit candidates ~4-5x more often than chance - and the ensemble wins overall.

method	NDCG@10	NDCG@50	MAP	P@10	vs random
random floor (shuffle)	0.049	0.084	0.061	0.047	1x
tfidf keyword match	0.205	0.235	0.137	0.189	4x
embeddings (MiniLM)	0.210	0.257	0.150	0.211	5x
ensemble (RRF) <- best overall	0.219	0.267	0.156	0.200	4x

Read it as: P@10 0.20 = ~2 of our top 10 are recruiter-confirmed good fits vs ~0.5 by chance (4x). Absolute numbers are modest because this is the hardest fair test: rank all 477 from scratch, count only strict Good-Fit, unjudged = miss.

- ▶ **Ensemble is best overall + 4-5x over random:** fusing three lenses generalises and is not overfit to one JD.
- ▶ **Rules/gate validated on the challenge data:** the off-domain 'Content Writer' drops from #1 to #21,804/100,000; 0 honeypots in the top 100.

And why each was chosen

pandas / numpy

vectorized featurization and scoring; no per-row loops, so 100k fits the budget.

scikit-learn

TF-IDF lexical recall, battle-tested and fast on CPU.

MiniLM (transformers)

compact offline sentence embeddings for semantic fit, precomputed once.

Reciprocal Rank Fusion

scale-free combination of heterogeneous lenses.

Docker

exact Stage-3 parity: CPU-only, offline reproduction in one command.

Google Colab

zero-setup hosted sandbox that runs end-to-end on a sample.

Everything needed to reproduce and verify

- ▶ **GitHub:** github.com/NavpreetDevpuri/canjob (public, real iterative history)
- ▶ **Reproduce:** `python rank.py --candidates ./candidates.jsonl --out ./submission.csv`
- ▶ **Sandbox:** Google Colab run-all notebook with a config panel and bundled sample
- ▶ **Docker:** `docker build -t canjob .` then `docker run --network none ...`
- ▶ **Web UI (bonus):** FastAPI + single-page app: upload JDs, manage candidates, run matches with live progress
- ▶ **Benchmark:** `python benchmark/run_benchmark.py` -> NDCG/MAP/P@10 on a public labelled dataset
- ▶ **Submission CSV:** top 100, UTF-8, candidate_id,rank,score,reasoning (validator-clean)
- ▶ **Metadata:** `submission_metadata.yaml` mirrors the portal fields
- ▶ **AI usage:** approach, architecture & decisions are the developer's; code written fast with Cursor; no LLM at rank time

Thank you

canjob - Intelligent Candidate Discovery & Ranking

Navpreet Singh | github.com/NavpreetDevpuri/canjob

Ideas & architecture by the developer; implemented quickly with Cursor (AI). No LLM/API at rank time.